

Министерство науки и высшего образования Российской Федерации
Федеральное государственное автономное
образовательное учреждение высшего образования
«НАЦИОНАЛЬНЫЙ ИССЛЕДОВАТЕЛЬСКИЙ УНИВЕРСИТЕТ ИТМО»

ОТЧЕТ по лабораторной работе
«Этап 3. Администрирование и оптимизация»
по дисциплине **«Проектирование баз данных»**

Автор:	Константинова Елизавета Анатольевна
Факультет:	ФИТиП
Группа:	М3203
Преподаватели:	Шевчик Софья Владимировна

ИТМО

Санкт-Петербург 2024

Запросы к бд:

Вывести всех пользователей и их роли на определенном сервере которые сейчас онлайн

```
SELECT u.username, ur.role_id
FROM servers AS s
LEFT JOIN users u ON u.user_id = s.owner_id
LEFT JOIN user_roles AS ur ON s.server_id = ur.server_id and u.user_id =
ur.user_id
WHERE u.online_status = TRUE and s.server_id = 97664
```

Вывести пары людей которые сидят как минимум на двух одинаковых серверах

```
SELECT DISTINCT
    u1.username AS username1,
    u2.username AS username2
FROM
    user_roles ur1
    JOIN user_roles ur2 ON ur1.server_id = ur2.server_id AND ur1.user_id
    <> ur2.user_id
    JOIN users u1 ON ur1.user_id = u1.user_id
    JOIN users u2 ON ur2.user_id = u2.user_id
GROUP BY
    u1.username, u2.username
HAVING
    COUNT(DISTINCT ur1.server_id) >= 2;
```

Вывести все сервера и их обладателей у которых есть хотя бы один текстовый канал

```
SELECT
    s.server_name,
    u.username AS owner_username
FROM
    servers s
    JOIN channels c ON s.server_id = c.server_id
    JOIN users u ON s.owner_id = u.user_id
WHERE
    c.channel_type = 'text_channel'
GROUP BY
    s.server_name, u.username;
```

Сервис, который вызывает Explain Analyze для каждого запроса и измеряет Cost:

Dockerfile:

```
FROM mcr.microsoft.com/dotnet/sdk:8.0 AS build-env
WORKDIR /discord

COPY . ./

RUN dotnet restore

RUN dotnet publish -c Release -o out
```

```
FROM mcr.microsoft.com/dotnet/runtime:8.0
WORKDIR /discord
COPY --from=build-env /discord/out .
EXPOSE 8080
ENTRYPOINT ["sh", "-c", "dotnet discord.dll"]
```

docker-compose:

```
analyze_app:
  restart: no
  container_name: analyze_app
  environment:
    ATTEMPT_COUNT: ${ATTEMPT_COUNT}
  build:
    context: ./discord
    dockerfile: Dockerfile
  volumes:
    - ./itogi:/docker-entrypoint-initdb.d/itogi
  depends_on:
    db:
      condition: service_healthy
  env_file:
    - .env
  ports:
    - "8080:8080"
```

Analyze.cs:

```
using System.Text.RegularExpressions;
using Npgsql;

namespace discord;

class Analyze
{
    static void Main(string[] args)
    {
        var timestamp = DateTime.Now.ToString("yyyyMMddHHmmss");
        var filePath = $"/docker-entrypoint-
initdb.d/itogi/query_itogi_{timestamp}.txt";
        var connectionString =
"Host=db;Username=postgres;Password=myspassword;Database=discord-
db;Port=5432;";
        Console.WriteLine("ANALYZE SERVICE STARTED");

        var attemptCount =
int.Parse(Environment.GetEnvironmentVariable("ATTEMPT_COUNT") ?? "4");

        string[] queries =
        {
            """
            SELECT u.username, ur.role_id
            FROM servers AS s
            LEFT JOIN users u ON u.user_id = s.owner_id
```

```

        LEFT JOIN user_roles AS ur ON s.server_id = ur.server_id and
u.user_id = ur.user_id
        WHERE u.online_status = TRUE and s.server_id = 97664
        """,
        """)
        SELECT DISTINCT
            u1.username AS username1,
            u2.username AS username2
        FROM
            user_roles ur1
            JOIN user_roles ur2 ON ur1.server_id = ur2.server_id AND
ur1.user_id <> ur2.user_id
            JOIN users u1 ON ur1.user_id = u1.user_id
            JOIN users u2 ON ur2.user_id = u2.user_id
        GROUP BY
            u1.username, u2.username
        HAVING
            COUNT(DISTINCT ur1.server_id) >= 2;
        """,
        """)
        SELECT
            s.server_name,
            u.username AS owner_username
        FROM
            servers s
            JOIN channels c ON s.server_id = c.server_id
            JOIN users u ON s.owner_id = u.user_id
        WHERE
            c.channel_type = 'text_channel'
        GROUP BY
            s.server_name, u.username;
        """)
    };

    var costs = new double[queries.Length][];
    for (var i = 0; i < costs.Length; i++)
    {
        costs[i] = new double[attemptCount];
    }

    try
    {
        using (var writer = new StreamWriter(filePath))
        {
            using (var connection = new
NpgsqlConnection(connectionString))
            {
                connection.Open();

                for (var q = 0; q < queries.Length; q++)
                {
                    var query = queries[q];
                    Console.WriteLine($"Executing query {q +
1}/{queries.Length}: {query}");
                    for (var a = 0; a < attemptCount; a++)
                    {
                        try

```

```

        {
            using (var command = new
NpgsqlCommand($"EXPLAIN ANALYZE {query}", connection))
            {
                using (var reader =
command.ExecuteReader())
                {
                    var res = new List<double>();

                    while (reader.Read())
                    {
                        var result = reader.GetString(0);
                        var match = Regex.Match(result,
@"cost=(\d+\.?\d+) .. (\d+\.?\d+)");

                        if (match.Success)
                        {
                            var startCost =
double.Parse(match.Groups[1].Value);
                            var endCost =
double.Parse(match.Groups[2].Value);

                            res.Add(endCost);
                        }
                    }

                    if (a == attemptCount - 1)
                    {
                        var minC = res.Min();
                        var maxC = res.Max();
                        var avgC = res.Average();
                        writer.WriteLine($"query:

{query}\n");

                        writer.WriteLine($"best case
time: {minC}");

                        writer.WriteLine($"worst case
time: {maxC}");

                        writer.WriteLine($"average case
time: {avgC}");

                        writer.WriteLine('\n');
                    }
                }
            }
        }
        catch (Exception ex)
        {
            Console.WriteLine($"ERROR: {ex.Message}");
        }
    }
}

Console.WriteLine($"RESULTS SAVED: {filePath}");
}
catch (Exception ex)
{
    Console.WriteLine($"error writing to file: {ex.Message}");
}

```

```

    }
}
}

```

Добавленные индексы:

```

CREATE INDEX idx_servers_owner_id ON servers (owner_id);
CREATE INDEX idx_user_roles_server_user_id ON user_roles (server_id,
user_id);
CREATE INDEX idx_user_roles_user_id_server_id ON user_roles (user_id,
server_id);
CREATE INDEX idx_channels_channel_type ON channels (channel_type);

```

Сравнительная таблица:

изначально	1 запрос	2 запрос	3 запрос
Best	8.3	1454.06	0.38
Worst	1814.93	25352528.95	3869.82
Avg	728.044	5058657.189999999	2428.444285714286
После индексов			
Best	0.88	1404.62	0.37
Worst	17.51	16744308.71	3682.29
Avg	10.324	4240299.414285715	1909.295
Лучше на	7180%	19,2%	27,2%

Создание партиции таблицы ролей пользователей:

```

ALTER TABLE user_roles RENAME TO user_roles_old;

CREATE TABLE user_roles
(
    user_role_id SERIAL PRIMARY KEY,
    user_id      INTEGER NOT NULL REFERENCES users (user_id),
    role_id      INTEGER NOT NULL REFERENCES roles (role_id),
    server_id    INTEGER NOT NULL REFERENCES servers (server_id)
) PARTITION BY RANGE (user_role_id);

ALTER TABLE user_roles_old
    ADD CONSTRAINT user_roles_old
        CHECK (user_role_id >= 1 AND user_role_id <= 100000);

CREATE TABLE user_roles_1 PARTITION OF user_roles
    FOR VALUES FROM
(
    1
) TO
(
    50000

```

```

);

CREATE TABLE user_roles_2 PARTITION OF user_roles
    FOR VALUES FROM
    (
        50000
    ) TO
    (
        100001
    );

WITH moved_rows AS (
DELETE
FROM user_roles_old a
WHERE user_role_id >= 1
    AND user_role_id < 50000 RETURNING a.*
)
INSERT
INTO user_roles_1
SELECT *
FROM moved_rows;

WITH moved_rows AS (
DELETE
FROM user_roles_old a
WHERE user_role_id >= 50000
    AND user_role_id < 100001 RETURNING a.*
)
INSERT
INTO user_roles_2
SELECT *
FROM moved_rows;

ALTER TABLE user_roles_old DROP CONSTRAINT user_roles_old;

```

изначально	1 запрос	2 запрос	3 запрос
Best	0.88	1404.62	0.37
Worst	17.51	16744308.71	3682.29
Avg	10.324	4240299.414285715	1909.295
После партиции			
Best	8.3	565.91	0.37
Worst	1820.78	16829892.86	3682.29
Avg	778.18	3313748.391	1909.295
Лучше на	-98%	30%	0%

Скрипт, который создает бэкапы базы данных каждые n-часов, последние m-бекапов хранятся

backup.sh:

```
#!/bin/bash

DB_USER=${DB_SUPERUSER}
DB_NAME=${POSTGRES_DB}
DB_PASSWORD=${POSTGRES_PASSWORD}

BACKUP_DIR="/backup"

BACKUP_COUNT=${M}

TIMESTAMP=$(date +%Y%m%d%H%M%S')
BACKUP_FILE="$BACKUP_DIR/backup_${TIMESTAMP}.sql"
sleep_time=${N}

export PGPASSWORD="$DB_PASSWORD"

mkdir -p "$BACKUP_DIR"

sleep "$sleep_time"

pg_dump -h db -p 5432 -U "$DB_USER" -d "$DB_NAME" > "$BACKUP_FILE"

BACKUP_FILES_COUNT=$(ls -1 "$BACKUP_DIR"/*.sql | wc -l)
if [ "$BACKUP_FILES_COUNT" -gt "$BACKUP_COUNT" ]; then
    ls -1t "$BACKUP_DIR"/*.sql | tail -n +$(( $BACKUP_COUNT + 1 )) | xargs rm -f
fi

unset PGPASSWORD

echo "backup"
```

Dockerfile:

```
FROM alpine:latest
RUN apk add --no-cache postgresql-client
COPY backup.sh ./backup.sh
RUN chmod +x /backup.sh
EXPOSE 8081

CMD ["sh", "/backup.sh"]
```

docker-compose:

```
backup:
  restart: always
  container_name: backup
```



```

environment:
  POSTGRES_DB: ${POSTGRES_DB}
  POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
  DB_USERNAME: ${DB_USERNAME}
  DB_SUPERUSER: ${DB_SUPERUSER}
  N: ${N}
  M: ${M}
build:
  context: ./backup
  dockerfile: Dockerfile
volumes:
  - ./backup:/backup
  - ./backup/backup.sh:/docker-entrypoint-initdb.d/backup.sh
depends_on:
  db:
    condition: service_healthy
env_file:
  - .env
ports:
  - "8081:8081"

```

.env:

```

POSTGRES_DB="discord-db"
POSTGRES_PASSWORD="mypassword"
DB_USERNAME="postgres"
DB_SUPERUSER="postgres"
FILLING_AMOUNT=100000
ATTEMPT_COUNT=10
N=10m
M=48

```

2 реплики Postgres с использованием Patroni:

```

docker-compose:
patroni:
  image: cybertec/pgwatch2
  container_name: patroni
  restart: unless-stopped
  depends_on:
    - db
  ports:
    - "2392:2392"
  environment:
    PATRONI_RESTAPI_CONNECT_ADDRESS: "0.0.0.0:2392"

replica1:
  image: postgres
  container_name: replica1
  restart: unless-stopped
  environment:
    POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
    POSTGRES_USER: ${DB_SUPERUSER}
    POSTGRES_DB: ${POSTGRES_DB}

```

```

    PGDATA: /var/lib/postgresql/data/pgdata
volumes:
  - replica1_data:/var/lib/postgresql/data
depends_on:
  - db

replica2:
  image: postgres
  container_name: replica2
  restart: unless-stopped
  environment:
    POSTGRES_PASSWORD: ${POSTGRES_PASSWORD}
    POSTGRES_USER: ${DB_SUPERUSER}
    POSTGRES_DB: ${POSTGRES_DB}
    PGDATA: /var/lib/postgresql/data/pgdata
  volumes:
    - replica2_data:/var/lib/postgresql/data
  depends_on:
    - db

volumes:
  replica2_data:
  replica1_data:

```

Итоговая структура проекта:

